

# Project: Investigate a Dataset - TMDb movie data

## Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
- [Conclusions](#)

## Introduction

### Dataset Description

Using the Movie Database(TMDb) with data from over 10,000 movies I will analyze the data to uncover trends among different movies and genres. This analysis will explore their popularity, runtimes, genres, and budget, as well as the correlations among these factors.

### Question(s) for Analysis

**Question 1:** What are the top 10 favorite movies?

**Question 2:** What year was the highest budget movie produce?

**Question 3:** What is the correlation between the average runtime of movies and the passage of time?

```
In [2]: # Import packates
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

# Data Wrangling

## General Properties

```
In [3]: #Load data
df = pd.read_csv('tmdb-movies.csv')

#view first 3 rows
df.head(3)
```

Out[3]:

	id	imdb_id	popularity	budget	revenue	original_title	cast	
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	http://www.

3 rows × 21 columns

```
In [4]: #show size of dataframe
df.shape
```

Out[4]: (10866, 21)

In [5]: *#dataset information*

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    10866 non-null  int64
1   imdb_id              10856 non-null  object
2   popularity            10866 non-null  float64
3   budget               10866 non-null  int64
4   revenue              10866 non-null  int64
5   original_title       10866 non-null  object
6   cast                 10790 non-null  object
7   homepage             2936 non-null   object
8   director             10822 non-null  object
9   tagline              8042 non-null   object
10  keywords             9373 non-null   object
11  overview             10862 non-null  object
12  runtime              10866 non-null  int64
13  genres               10843 non-null  object
14  production_companies  9836 non-null   object
15  release_date         10866 non-null  object
16  vote_count           10866 non-null  int64
17  vote_average         10866 non-null  float64
18  release_year         10866 non-null  int64
19  budget_adj           10866 non-null  float64
20  revenue_adj          10866 non-null  float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

In [6]: *#data summary*

```
df.describe()
```

Out[6]:

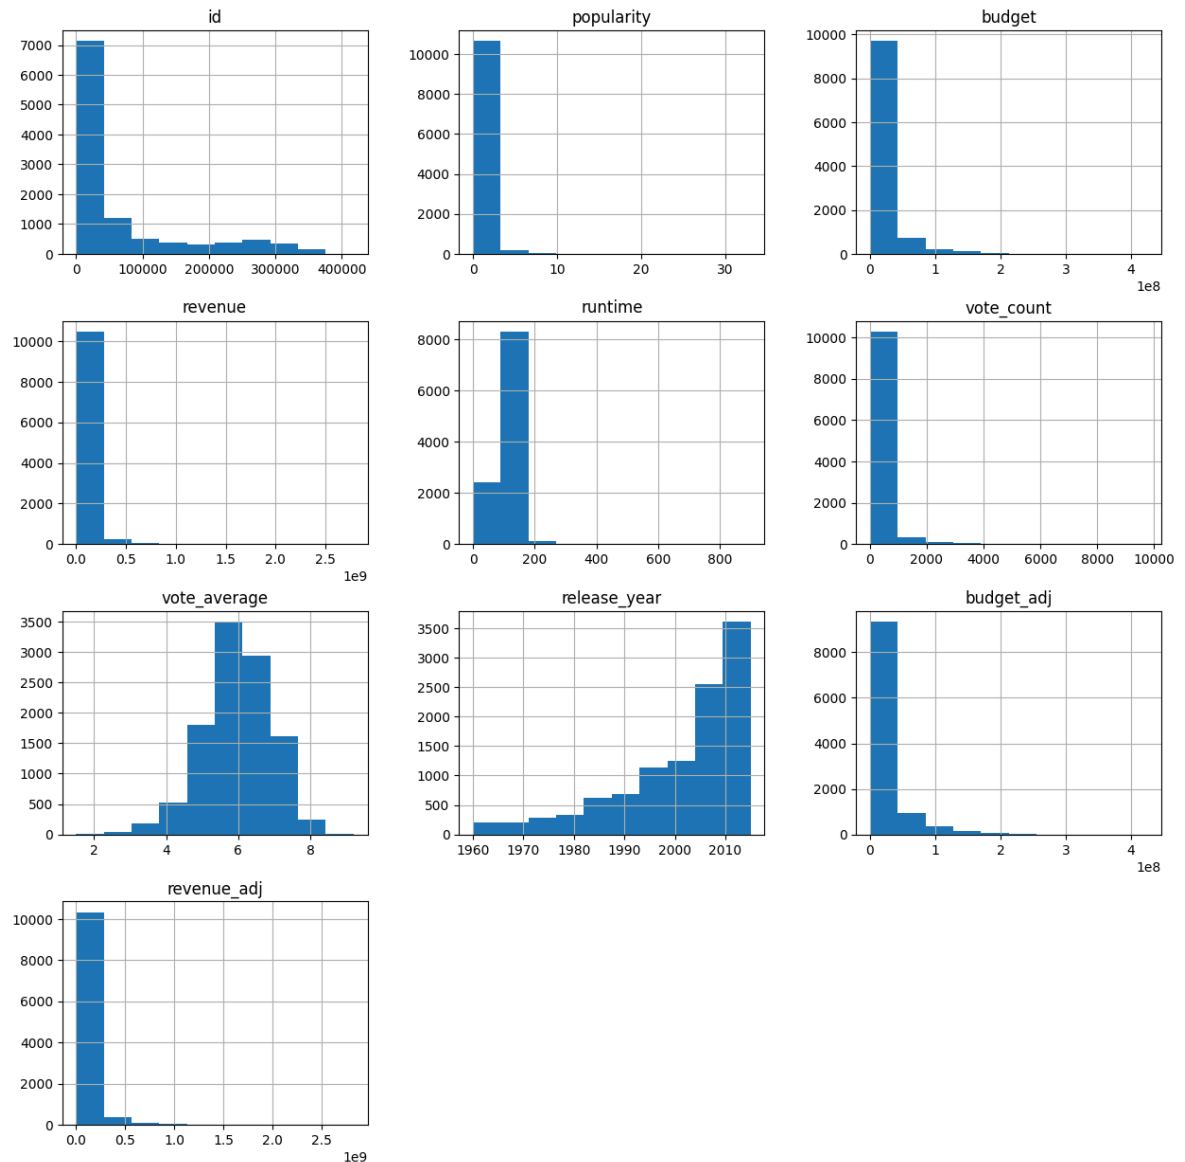
	id	popularity	budget	revenue	runtime	vote_count	v
<b>count</b>	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000	10866.000000	1
<b>mean</b>	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863	217.389748	
<b>std</b>	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405	575.619058	
<b>min</b>	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000	10.000000	
<b>25%</b>	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000	17.000000	
<b>50%</b>	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000	38.000000	
<b>75%</b>	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000	145.750000	
<b>max</b>	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	

```
In [7]: #show data types
df.nunique()
```

```
Out[7]: id                10865
imdb_id                 10855
popularity             10814
budget                  557
revenue                 4702
original_title         10571
cast                   10719
homepage                2896
director                5067
tagline                 7997
keywords                8804
overview               10847
runtime                 247
genres                  2039
production_companies    7445
release_date            5909
vote_count              1289
vote_average            72
release_year            56
budget_adj              2614
revenue_adj             4840
dtype: int64
```

```
In [8]: df.hist(figsize=(15,15))
```

```
Out[8]: array([[<AxesSubplot: title={'center': 'id'}>,
  <AxesSubplot: title={'center': 'popularity'}>,
  <AxesSubplot: title={'center': 'budget'}>],
  [<AxesSubplot: title={'center': 'revenue'}>,
  <AxesSubplot: title={'center': 'runtime'}>,
  <AxesSubplot: title={'center': 'vote_count'}>],
  [<AxesSubplot: title={'center': 'vote_average'}>,
  <AxesSubplot: title={'center': 'release_year'}>,
  <AxesSubplot: title={'center': 'budget_adj'}>],
  [<AxesSubplot: title={'center': 'revenue_adj'}>, <AxesSubplot: >,
  <AxesSubplot: >]], dtype=object)
```



## Data Cleaning

```
In [9]: #check for duplicates
df.duplicated().sum()
```

Out[9]: 1

```
In [10]: #Drop duplicate rows found
df.drop_duplicates(inplace=True)

#check to make sure duplicates have been dropped
df.duplicated().sum()
```

Out[10]: 0

```
In [11]: #Look for missing/NULL values
df.isna().sum()
```

```
Out[11]: id                0
imdb_id                  10
popularity               0
budget                  0
revenue                  0
original_title           0
cast                    76
homepage                7929
director                 44
tagline                 2824
keywords                1493
overview                 4
runtime                  0
genres                   23
production_companies    1030
release_date             0
vote_count               0
vote_average             0
release_year             0
budget_adj               0
revenue_adj              0
dtype: int64
```

```
In [16]: #drop columns that are not necessary
df.drop(columns=['imdb_id', 'homepage', 'director', 'tagline', 'keywords', 'overview'])
```

```
Out[16]:
```

	id	popularity	budget	revenue	original_title	cast	runtime	
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	124	Ac
1	76341	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	120	Ac
2	262500	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	119	
3	140607	11.173104	200000000	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D...	136	Ac
4	168259	9.335014	190000000	1506249360	Furious 7	Vin Diesel Paul Walker Jason Statham Michelle ...	137	
...	...	...	...	...	...	...	...	...
10861	21	0.080598	0	0	The Endless Summer	Michael Hynson Robert August Lord 'Tally Ho' B...	95	
10862	20379	0.065543	0	0	Grand Prix	James Garner Eva Marie Saint Yves Montand Tosh...	176	A
10863	39768	0.065141	0	0	Beregis Avtomobilya	Innokentiy Smoktunovskiy Oleg Efremov Georgi Z...	94	
10864	21449	0.064317	0	0	What's Up, Tiger Lily?	Tatsuya Mihashi Akiko Wakabayashi Mie Hama Joh...	80	
10865	22293	0.035919	19000	0	Manos: The Hands of Fate	Harold P. Warren Tom Neyman John Reynolds Dian...	74	

10865 rows × 12 columns

```
In [12]: print(df.columns)

Index(['id', 'imdb_id', 'popularity', 'budget', 'revenue', 'original_title',
      'cast', 'homepage', 'director', 'tagline', 'keywords', 'overview',
      'runtime', 'genres', 'production_companies', 'release_date',
      'vote_count', 'vote_average', 'release_year', 'budget_adj',
      'revenue_adj'],
      dtype='object')
```

```
In [ ]: #Remove rows that have '0' in Budget and revenue
```

## Exploratory Data Analysis

### Research Question 1:

What are the top 10 favorite movies?

```
In [13]: #top 10 popular movies
top_10_movies = df.sort_values(by='popularity', ascending=False).head(10)

# Display the top 10 favorite movies
print(top_10_movies[['original_title', 'popularity']])
```

	original_title	popularity
0	Jurassic World	32.985763
1	Mad Max: Fury Road	28.419936
629	Interstellar	24.949134
630	Guardians of the Galaxy	14.311205
2	Insurgent	13.112507
631	Captain America: The Winter Soldier	12.971027
1329	Star Wars	12.037933
632	John Wick	11.422751
3	Star Wars: The Force Awakens	11.173104
633	The Hunger Games: Mockingjay - Part 1	10.739009

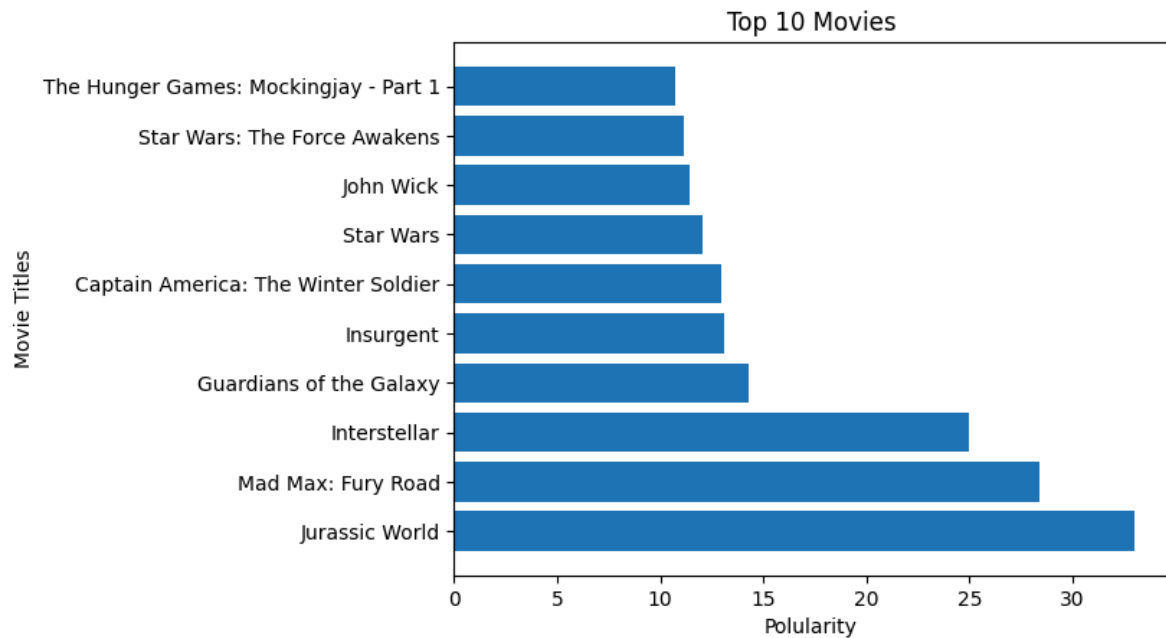


```
In [15]: #plot results
plt.barh(top_10_movies['original_title'], top_10_movies['popularity']);

#set title
plt.title('Top 10 Movies')

#set axis titles
plt.xlabel('Popularity')
plt.ylabel('Movie Titles')
```

```
Out[15]: Text(0, 0.5, 'Movie Titles')
```



From the data shown you can find that Jurassic world is the leading top favorite movie

## Research Question 2:

What year was the highest budget movie produced?

```
In [19]: #top 10 popular movies
highest_budget = df.sort_values(by='budget', ascending=False).head(15)

# Display the top 10 favorite movies
print(highest_budget[['original_title', 'release_year', 'budget']])
```

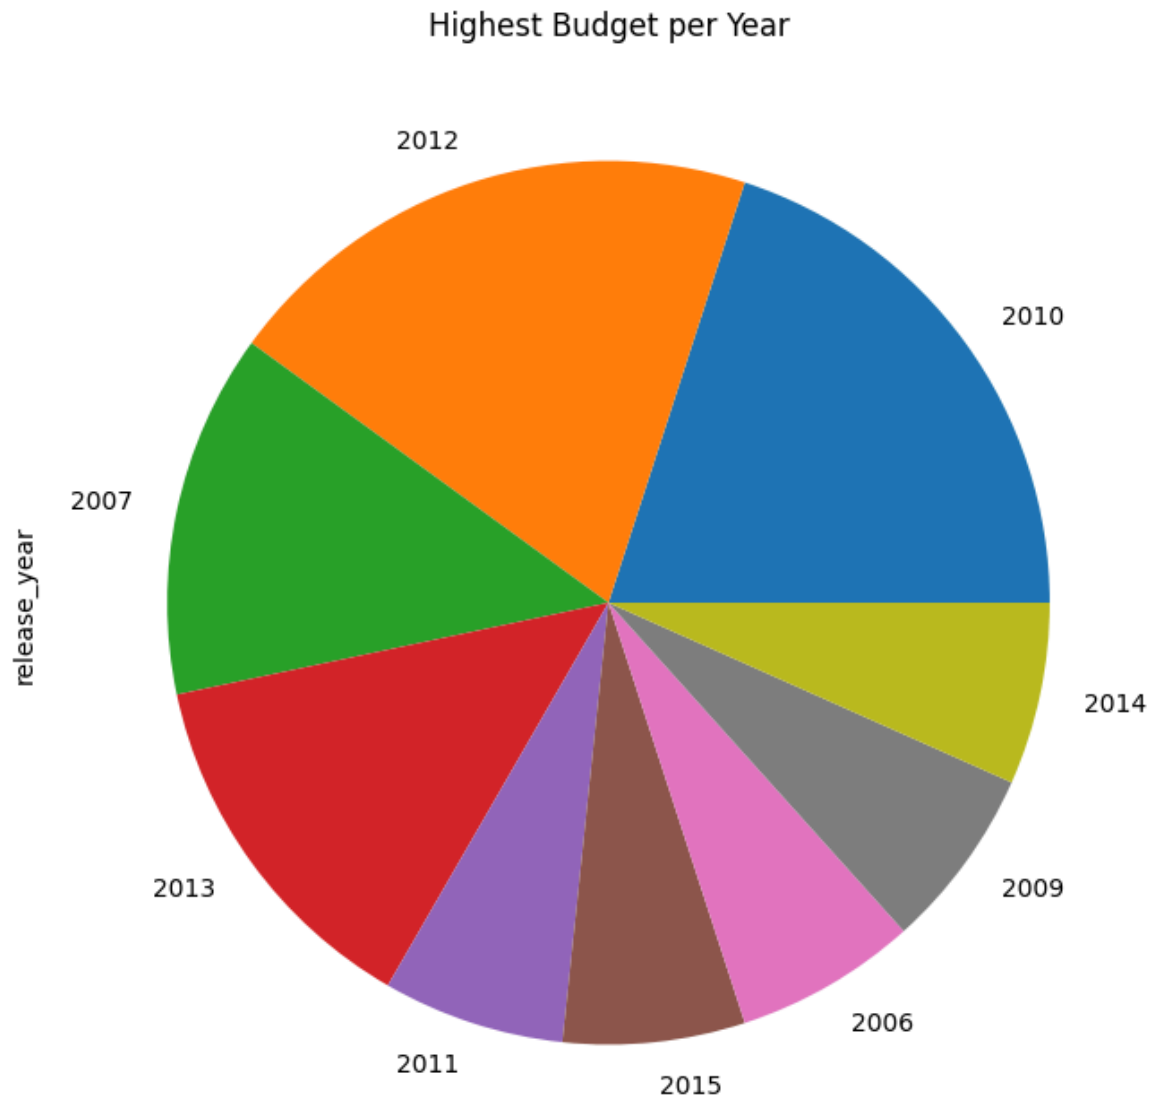
	original_title	release_year	budget
2244	The Warrior's Way	2010	425000000
3375	Pirates of the Caribbean: On Stranger Tides	2011	380000000
7387	Pirates of the Caribbean: At World's End	2007	300000000
14	Avengers: Age of Ultron	2015	280000000
6570	Superman Returns	2006	270000000
4411	John Carter	2012	260000000
1929	Tangled	2010	260000000
7394	Spider-Man 3	2007	258000000
5508	The Lone Ranger	2013	255000000
4367	The Hobbit: An Unexpected Journey	2012	250000000
1923	Harry Potter and the Deathly Hallows: Part 1	2010	250000000
1389	Harry Potter and the Half-Blood Prince	2009	250000000
5431	The Hobbit: The Desolation of Smaug	2013	250000000
643	X-Men: Days of Future Past	2014	250000000
4363	The Dark Knight Rises	2012	250000000

```
In [20]: highest_budget_year = highest_budget.loc[highest_budget['budget'].idxmax(), 'release_year']
print(f'Highest Budget Year is: {highest_budget_year}')
```

Highest Budget Year is: 2010

```
In [78]: #plot years using pie chart to visually see what year had the highest budget
highest_budget['budget'],highest_budget['release_year'].value_counts().plot(kind='pie',
print(r'))
```

Out[78]: Text(0.5, 1.0, 'Highest Budget per Year')



This chart shows a clear understanding of which year the highest budget was produced but after I made this Pie chart I realized without original\_title names it is difficult to tell which movie had the highest budget per year made.

I tried it another way to include a legend making the Pie chart easier to understand quickly while simultaneously revising this project to include a user-defined function as requested.

```

In [26]: # Plot highest budget movie by release year with movie titles as legend
def plot_highest_budget_pie_chart(df):

    # Find the movie with the highest budget per year
    highest_budget_per_year = df.loc[df.groupby('release_year')['budget'].idxmax()]


    # Create the pie chart for the highest budget by year
    fig, ax = plt.subplots(figsize=(8, 8))
    ax.pie(highest_budget_per_year['budget'], labels=highest_budget_per_year['release_year'],
           autopct='%1.1f%%')

    # Add a Legend with the highest budget movie title for each year
    ax.legend(highest_budget_per_year['original_title'], title="Movie Titles",
              bbox_to_anchor=(1.05, 1), loc='best')

    # Display the chart
    plt.title('Highest Budget Movie Per Year')
    plt.show()

# Call the function
plot_highest_budget_pie_chart(highest_budget)

```



From analyzing the data we can find that in 2010 the movie 'The Warrior's Way' had the largest budget in the amount of 425,000,000 USD which is almost 100,000,000 USD more than the next highest budget movie made.

### Research Question 3:

What is the correlation between the average runtime of movies and the passage of time?

```
In [81]: #find the average movie runtime  
df['runtime'].mean()
```

```
Out[81]: 102.07086324314375
```

```
In [87]: #find the minimum movie runtime  
df['runtime'].min()
```

```
Out[87]: 0
```

```
In [88]: #find the maximum movie runtime  
df['runtime'].max()
```


```
Out[88]: 900
```

```
In [86]: #Group by 'release_year' and calculate the average 'runtime'
average_runtime_per_year=df.groupby('release_year')['runtime'].mean()

#Plot a line chart
plt.plot(average_runtime_per_year.index, average_runtime_per_year.values)

#Add labels ad title
plt.xlabel('Release Year')
plt.ylabel('Average Runtime (Minutes)')
plt.title('Average Movie runtime over the years')

#Show the plot
plt.grid(True)
plt.show()
```



## Conclusions

**Question 1: What are the top 10 favorite movies?**

Jurassic world is the top favorite movie of all times.

**Question 2: What year was the highest budget movie produce?**

The Warrior's Way was the movie that had the largest budget and was made in 2010

### Question 3: What is the correlation between the average runtime of movies and the passage of time?

The included graph of movie runtimes shows a correlation between the length of movies and time, indicating that movies are getting shorter over the years.

#### Side Note:

I found the material for this project quite challenging and it took me a considerable amount of time to complete. I wish there had been a step-by-step video guide on getting started and navigating the workspaces before beginning. Loading the dataframe at the start was particularly difficult. However, after spending many hours experimenting in the workspaces, I began to enjoy the process and ultimately feel that I learned a lot.

## Limitations

After I reviewed the data a second time I realized that I should probably check the min and max runtimes when trying to find a correlation between the average runtime of movies and time. This is where I realized that some of the data may in fact be short clips, with little to no time. Or a long TV series with extra long run times which lead me to believe I may not have an accurate

```
In [27]: # Running this cell will execute a bash command to convert this notebook to an
!python -m nbconvert --to html Investigate_a_Dataset.ipynb
```

```
[NbConvertApp] WARNING | pattern 'Investigate_a_Dataset_revised.ipynb' matched no files
```

```
This application is used to convert notebook files (*.ipynb)
to various other formats.
```

```
WARNING: THE COMMANDLINE INTERFACE MAY CHANGE IN FUTURE RELEASES.
```

#### Options

```
=====
```

```
The options below are convenience aliases to configurable class-options,
as listed in the "Equivalent to" description-line of the aliases.
```

```
To see all configurable class-options for some <cmd>, use:
```

```
<cmd> --help-all
```

#### --debug

```
set log level to logging.DEBUG (maximize logging output)
```

```
Equivalent to: [--Application.log_level=10]
```

#### --show-config

```
Show the application's configuration (human-readable format)
```

In [ ]: